Containers vs Virtual Machines

Containers and Virtual Machines (VMs) offer different ways to deploy applications to a web server. While VMs share only the hardware of the host system, Containers share both the hardware and the OS. To compare both approaches, I created a simple web application for Contact Management System using HTML, CSS, and JS for front-end and Python Flask and SQLite for back-end.

I successfully deployed the application on VMs using Vagrant and VirtualBox, and I did the same on Containers using Docker. One of the superficial differences observed between the two is that while the codebase remains the same, we are changing the supporting files for each use case. For deploying to VM, I required a **VagrantFile**, whereas for deploying to Containers, I used **Dockerfile**. Both these files essentially perform similar jobs which is to define & set up the environment, and deploy the application using the given commands.

However, to statistically define which approach is better, we want to observe how their performancse over various metrics such as:

1. Startup time - The time it takes for an application to become fully operational from the moment it is started.
2. Memory usage - The amount of system memory (RAM) an application consumes while running.
3. CPU utilization - The percentage of CPU resources used by an application relative to the total available CPU resources.
4. Request throughput - The number of requests an application can process in a given time frame, usually measured in requests per second (RPS).
5. Response time - The amount of time taken by an application to respond to a request, typically measured from the moment a request is received until a response is sent back.
6. Garbage collection - The process of automatically freeing up memory by reclaiming space occupied by objects that are no longer in use by the application.

Startup Time

For containers, I wrote a shell script which encapsulated image creation and containerization commands and returned the startup time as soon as the application was deployed to the container. On the other hand, I manually calculated using a timer on my system for VMs. I observed that while the containers were up within 15-20 seconds, the VMs would take upto 4-5 minutes for the same. With respect to startup time, Containers have a clear edge over VMs.

The first image below represents the startup time for **Containers (17 secs)**, whereas the second image represents the same for **VMs (3 minute, 57 seconds)**.

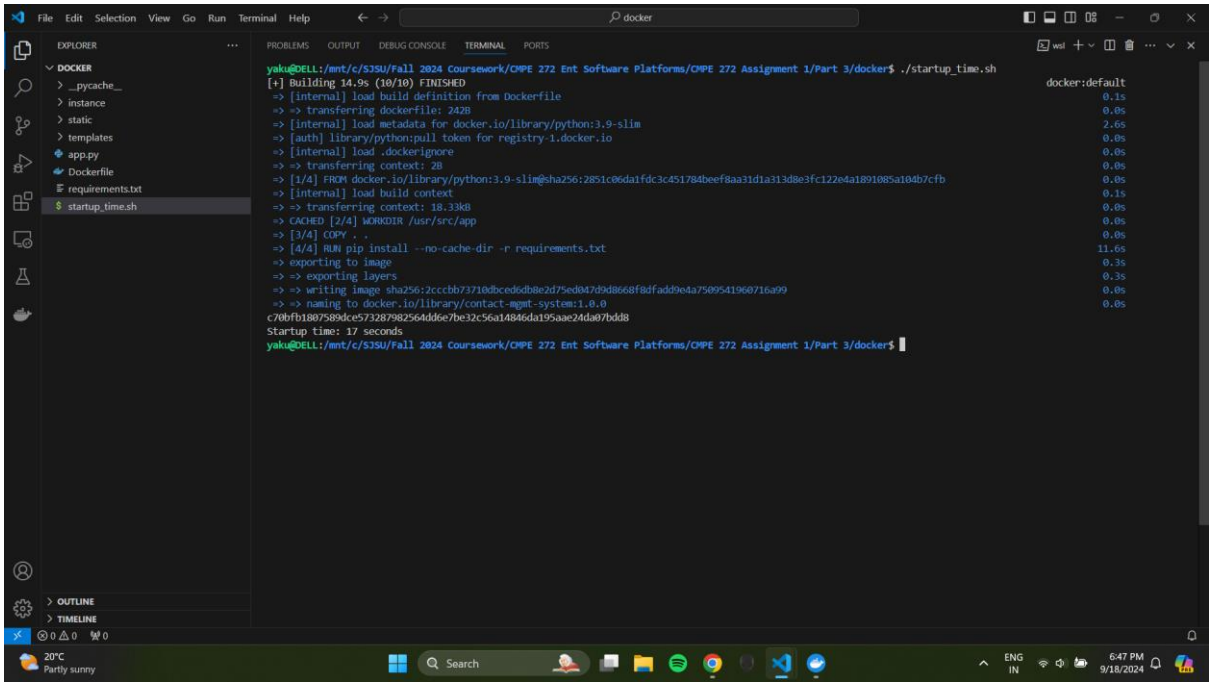

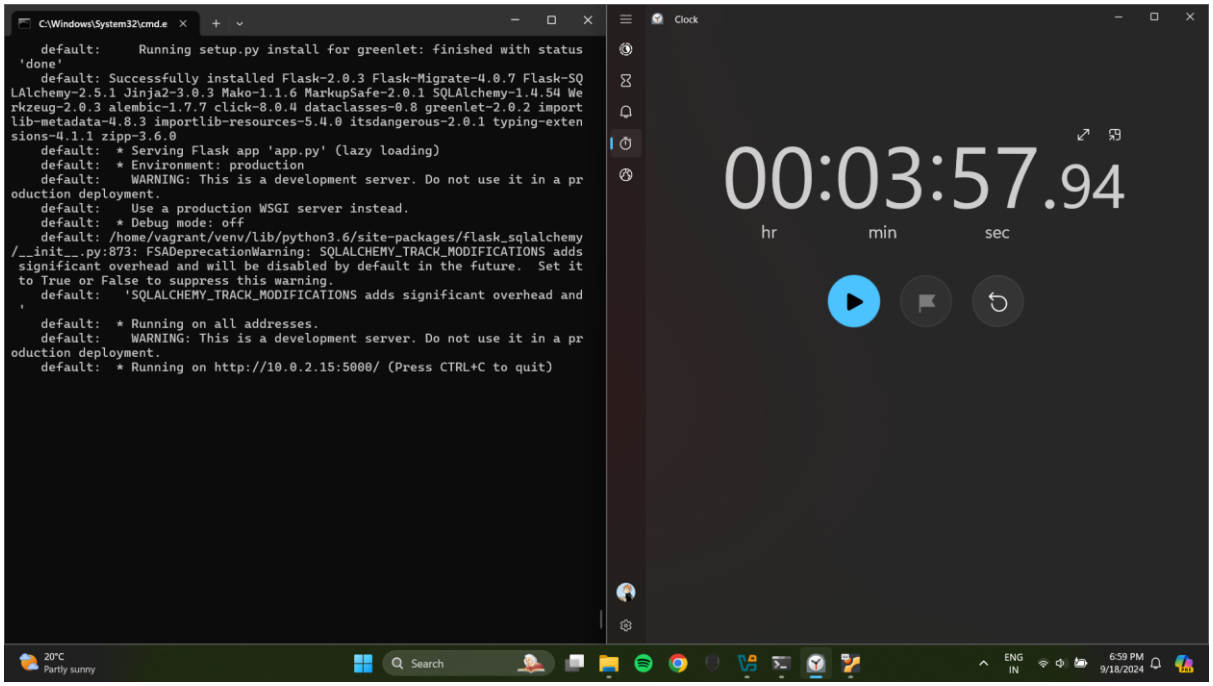

Memory Usage & CPU Utilization

With respect to memory usage & cpu utilization, I used commands for getting the stats for both. For dockers, I was able to do so using the **docker stats** command. On the other hand, for VMs, I used the **ps** command. With respect to memory usage, I observed that the containers used significantly lesser memory for the same functionality. Where VMs used nearly 5% of the memory, Containers used closer to 1%. On the other hand, regarding CPU Utilization, it was a lot closer with Containers using 0.01% and VMs using 0.00%. VMs did use much lesser, but neither is using a lot of CPU.

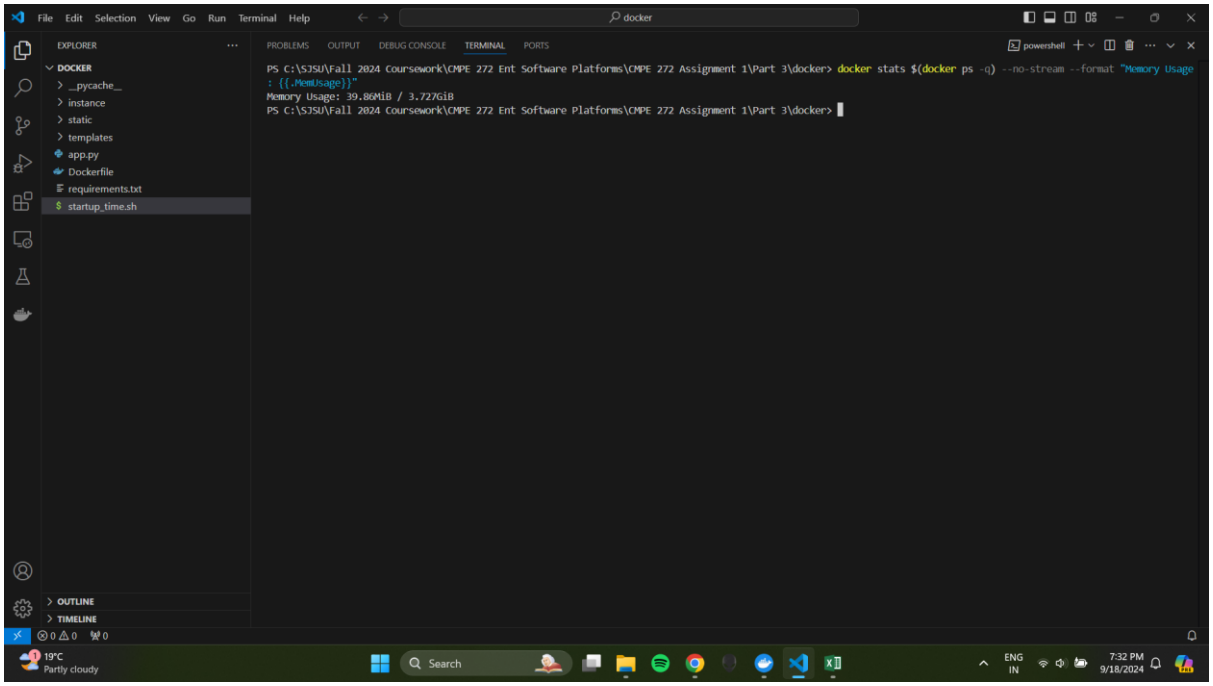
File Edit Selection View Go Run Terminal Help
docker
EXPLORER
DOCKER
__pycache__
instance
static
templates
app.py
Dockerfile
requirements.txt
startup_time.sh
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
powershell
PS C:\SJSU\Fall 2024 Coursework\CMPE 272 Ent Software Platforms\CMPE 272 Assignment 1\Part 3\docker> docker stats $(docker ps -q) --no-stream --format "Memory Usage: {{.MemUsage}}"
Memory Usage: 39.86MiB / 3.727GiB
PS C:\SJSU\Fall 2024 Coursework\CMPE 272 Ent Software Platforms\CMPE 272 Assignment 1\Part 3\docker>
OUTLINE
TIMELINE
19°C
Partly cloudy
Search
ENG
IN
7:32 PM
9/18/2024

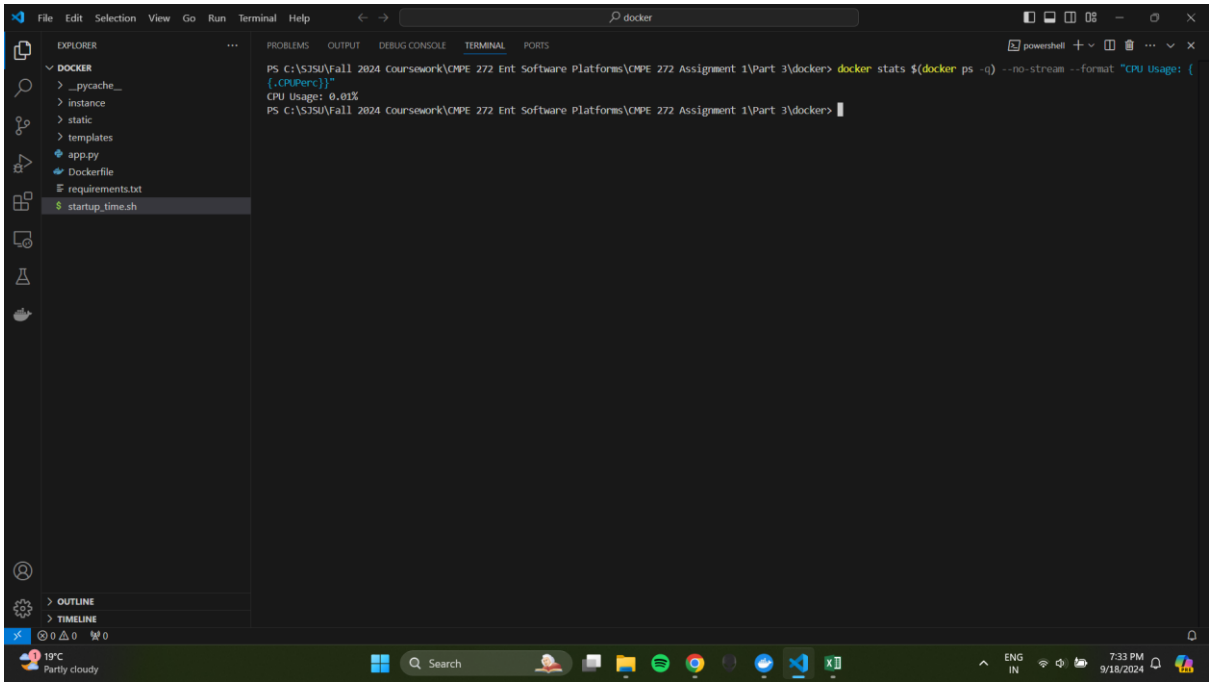
File Edit Selection View Go Run Terminal Help
docker
EXPLORER
DOCKER
__pycache__
instance
static
templates
app.py
Dockerfile
requirements.txt
startup_time.sh
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
powershell
PS C:\SJSU\Fall 2024 Coursework\CMPE 272 Ent Software Platforms\CMPE 272 Assignment 1\Part 3\docker> docker stats $(docker ps -q) --no-stream --format "CPU Usage: {{.CPUPerc}}"
CPU Usage: 0.01%
PS C:\SJSU\Fall 2024 Coursework\CMPE 272 Ent Software Platforms\CMPE 272 Assignment 1\Part 3\docker>
OUTLINE
TIMELINE
19°C
Partly cloudy
Search
ENG
IN
7:33 PM
9/18/2024

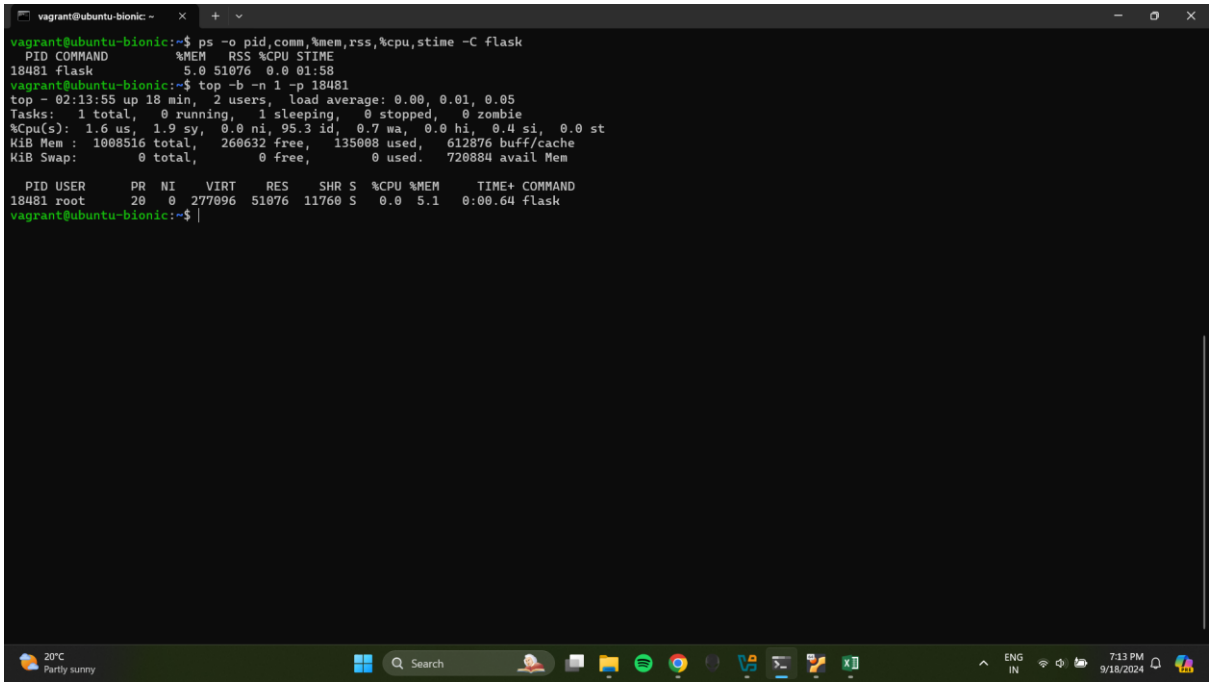

The first two images show the Memory and CPU usage in the Containers approach, whereas the third image shows the same for VMs. The table below summarizes the observed data as well.

Metric	VMs	Containers
Memory Usage %	5.10%	1.04%
CPU Utilization %	0.00%	0.01%

Request Throughput & Response Time

Request throughput and Response time go hand-in-hand, and hence, I used the Apache Benchmarking tool to measure both these metrics. I tested both VMs and Containers in two different scenarios:

a. Firing 100 requests individually, one at a time.
b. Firing 100 requests concurrently, in a batch size of 10.

Surprisingly, I observed that for the first scenario, VMs were quicker than Containers to handle individual requests. On the other hand, Containers fared better with the second scenario. The metrics are represented in the table below.

Scenario	Metric	VMs	Containers
Individual	Request Throughput (Requests/sec)	217.63	78.78
	Median Response Time (ms)	4	12
	Max Response Time (ms)	7	17
Concurrent	Request Throughput (Requests/sec)	240.90	430.82
	Median Response Time (ms)	40	21
	Max Response Time (ms)	57	42

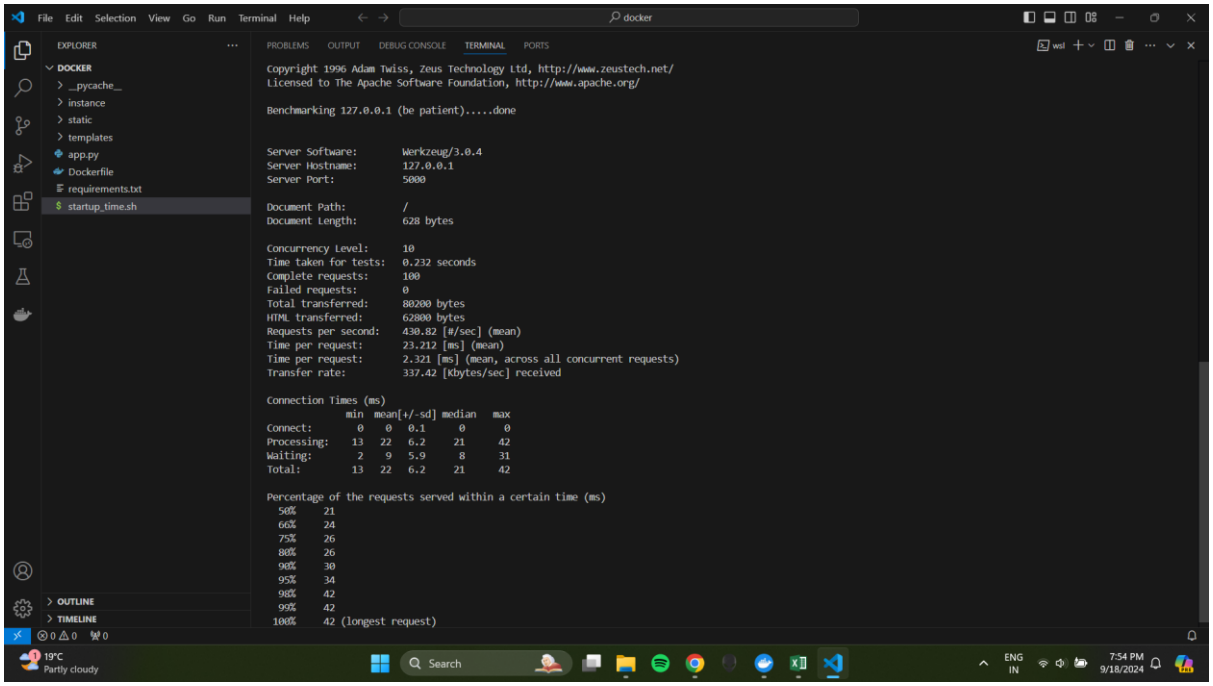

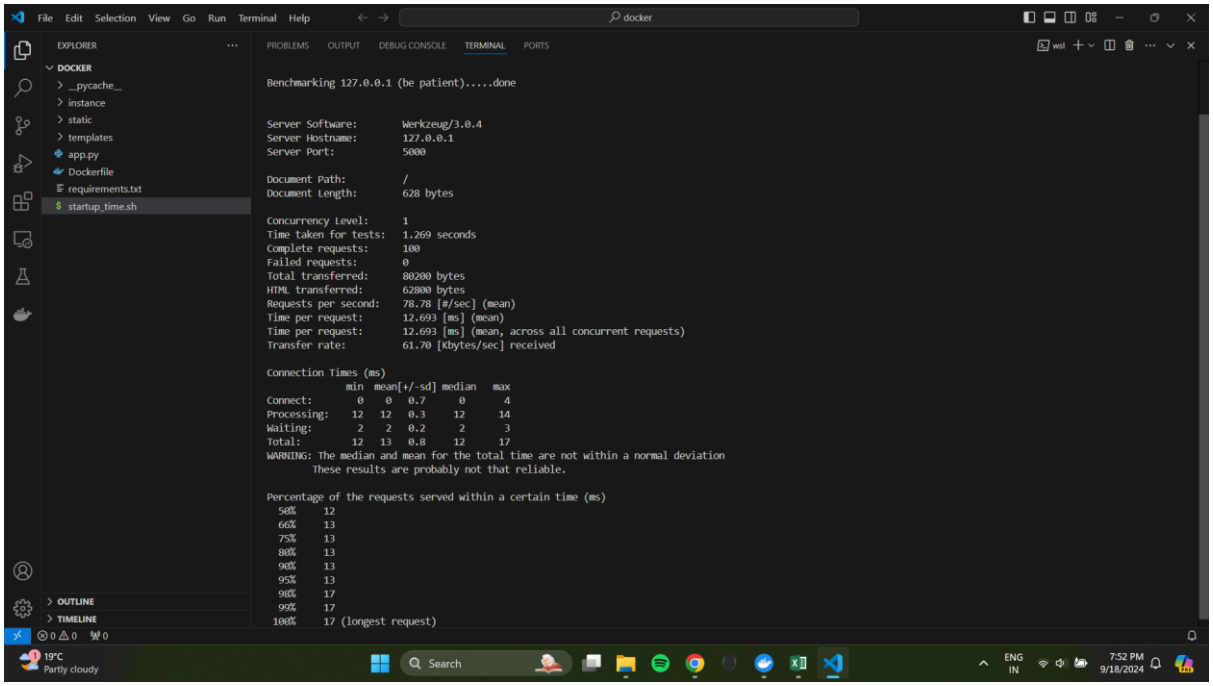

These two images show the Apache Benchmark results for individual and concurrent requests scenarios in a Container environment respectively.

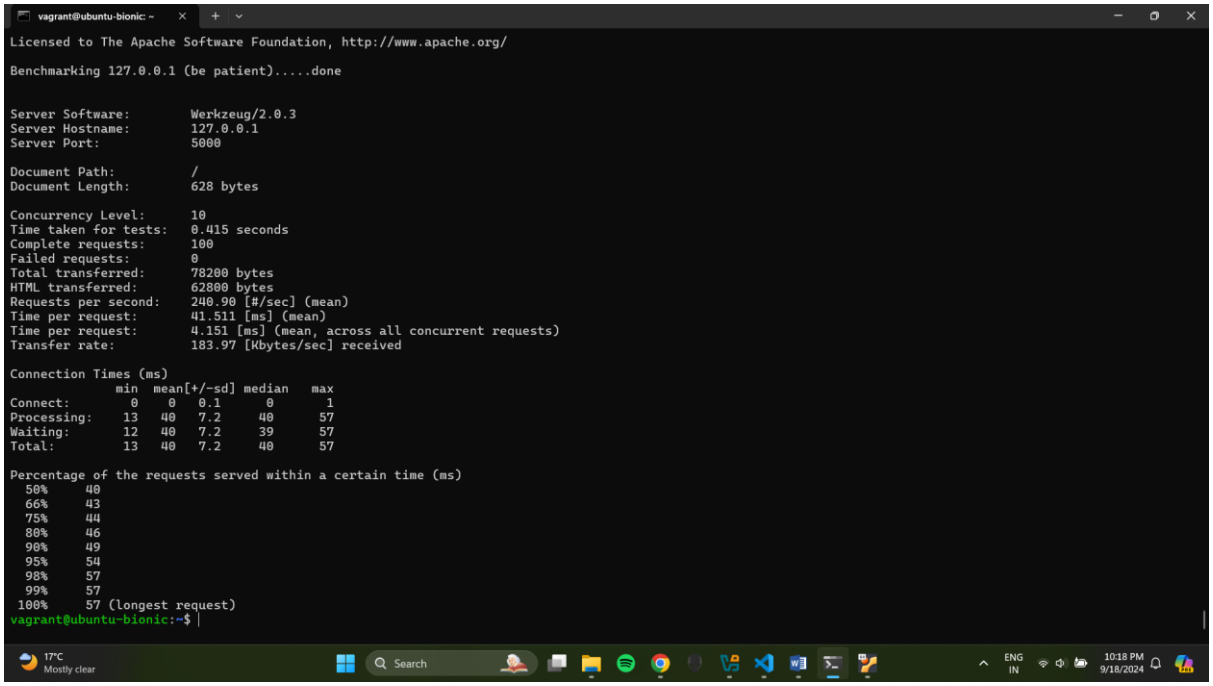

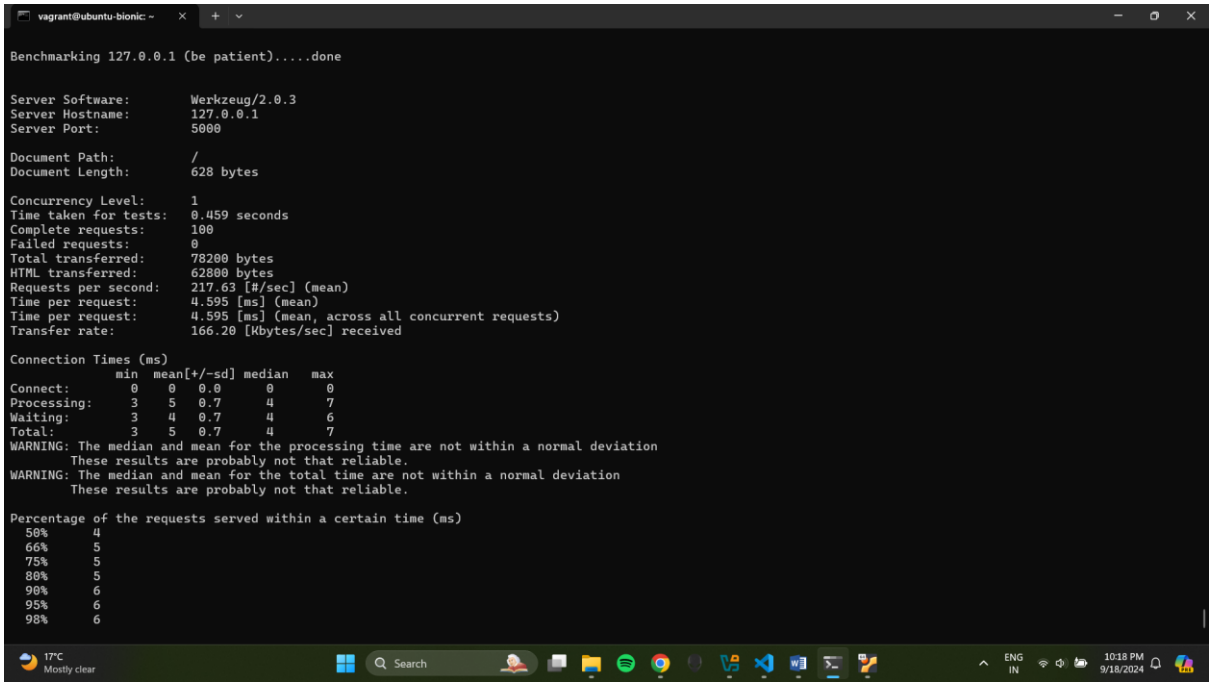

These two images show the Apache Benchmark results for individual and concurrent requests scenarios in a VM environment respectively.

Garbage Collection

To track Garbage Collection, I implemented a unique route (‘/gc_stats’) from my main webpage to access the Garbage Collection stats. I implemented this using Flask (with GC import) the same way I implemented the rest of the application. To keep it as similar as possible, I performed the same set of tasks in a row (add contact, edit, delete) before checking garbage collection page. The difference observed isn’t major but the VM did perform higher collections than the Container.

The results for the same are displayed below in the two images for the Container and VM respectively.

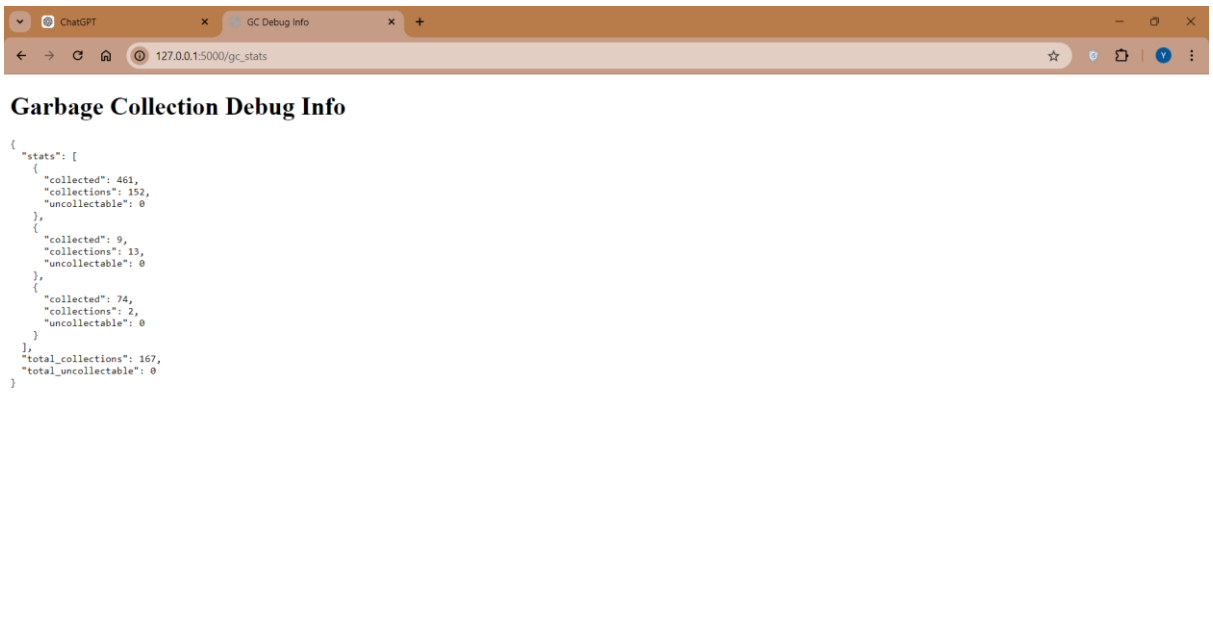

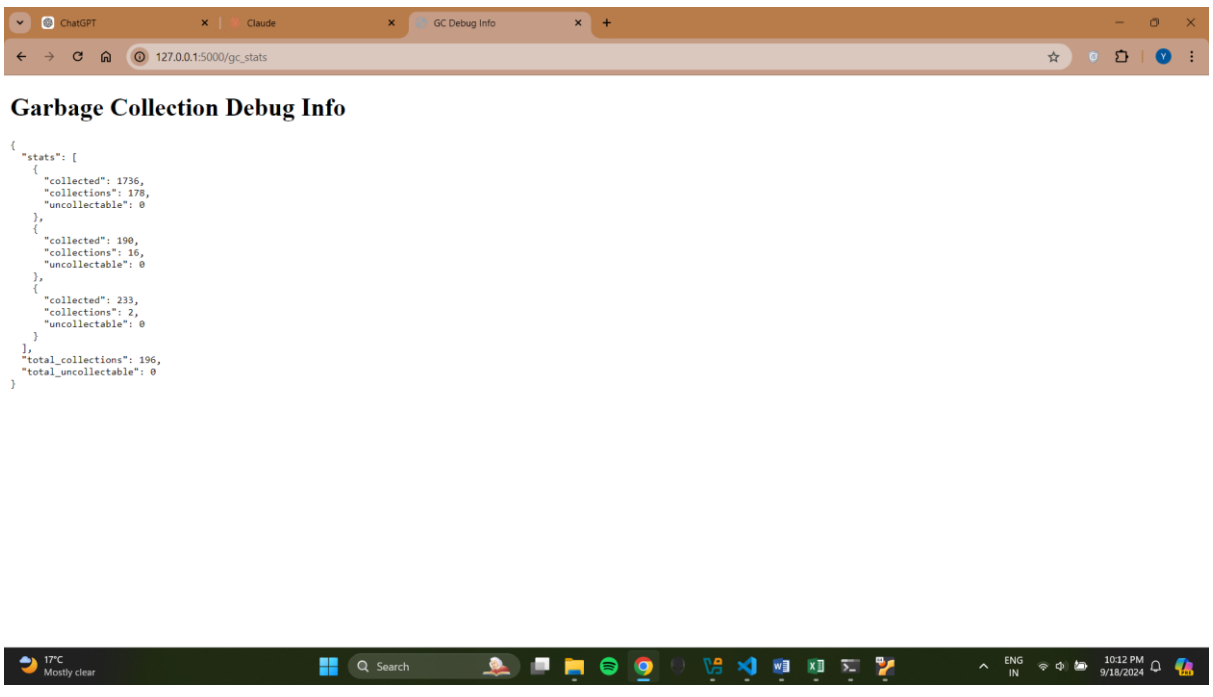

Summary

To summarize my findings, while there are certainly scenarios where the metrics favoured VMs, they are overwhelmingly favourable to Containers, especially in terms of Startup Time and Memory Usage, both of which are quite important. While there might still have use cases that prefer VMs, Containers are largely better thanks to the extra OS abstraction and performance.